

CmpE 462 : Machine Learning

Assignment 2

Name & Surname : Emre Girgin

ID : 2016400099

Part 1

In this part of the assignment I've trained 6 different Logistic Regression models. (2 batch sizes x 3 Learning rates) Only van and saab classes are taken. The features of the samples are normalized based on **MinMaxNormalization**. For stopping the training, I've checked the training loss of the last two epochs. If their difference is lower than a threshold, training stops. As the learning rate I've tested with 0.01, 0.5 and 1.0. Note that the x-axis of the figures is the iterations, not the epochs.

If you want to see additional figures and the training details, make sure that you have made *debug* variable from the code is **True**.

Iteration: 470	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5982	Val Loss: 0.7037	Accuracy: 0.59	Diff: 0.00010
Iteration: 471	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5981	Val Loss: 0.7036	Accuracy: 0.59	Diff: 0.00010
Iteration: 472	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5980	Val Loss: 0.7035	Accuracy: 0.59	Diff: 0.00010
Iteration: 473	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5979	Val Loss: 0.7034	Accuracy: 0.59	Diff: 0.00010
Iteration: 474	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5978	Val Loss: 0.7033	Accuracy: 0.59	Diff: 0.00010
Iteration: 475	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5977	Val Loss: 0.7032	Accuracy: 0.59	Diff: 0.00010
Iteration: 476	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5976	Val Loss: 0.7031	Accuracy: 0.59	Diff: 0.00010
Iteration: 477	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5975	Val Loss: 0.7030	Accuracy: 0.59	Diff: 0.00010
Iteration: 478	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5974	Val Loss: 0.7030	Accuracy: 0.59	Diff: 0.00010
Iteration: 479	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5973	Val Loss: 0.7029	Accuracy: 0.59	Diff: 0.00010
Iteration: 480	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5972	Val Loss: 0.7028	Accuracy: 0.59	Diff: 0.00010
Iteration: 481	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5971	Val Loss: 0.7027	Accuracy: 0.59	Diff: 0.00010
Iteration: 482	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5970	Val Loss: 0.7026	Accuracy: 0.59	Diff: 0.00010
Iteration: 483	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5969	Val Loss: 0.7025	Accuracy: 0.59	Diff: 0.00010
Iteration: 484	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5968	Val Loss: 0.7024	Accuracy: 0.59	Diff: 0.00010
Iteration: 485	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5967	Val Loss: 0.7023	Accuracy: 0.59	Diff: 0.00010
Iteration: 486	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5966	Val Loss: 0.7022	Accuracy: 0.59	Diff: 0.00010
Iteration: 487	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5965	Val Loss: 0.7021	Accuracy: 0.59	Diff: 0.00010
Iteration: 488	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5964	Val Loss: 0.7020	Accuracy: 0.59	Diff: 0.00010
Iteration: 489	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5963	Val Loss: 0.7019	Accuracy: 0.59	Diff: 0.00010
Iteration: 490	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5962	Val Loss: 0.7018	Accuracy: 0.60	Diff: 0.00010
Iteration: 491	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5961	Val Loss: 0.7017	Accuracy: 0.60	Diff: 0.00010
Iteration: 492	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5960	Val Loss: 0.7016	Accuracy: 0.60	Diff: 0.00010
Iteration: 493	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5959	Val Loss: 0.7015	Accuracy: 0.60	Diff: 0.00010
Iteration: 494	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5958	Val Loss: 0.7014	Accuracy: 0.60	Diff: 0.00010
Iteration: 495	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5957	Val Loss: 0.7013	Accuracy: 0.60	Diff: 0.00010
Iteration: 496	FoldID: 1	LearningRate: 0.01	Batch Size: 332	Train Loss: 0.5956	Val Loss: 0.7012	Accuracy: 0.60	Diff: 0.00010

Training details when debug is True.

Step 1

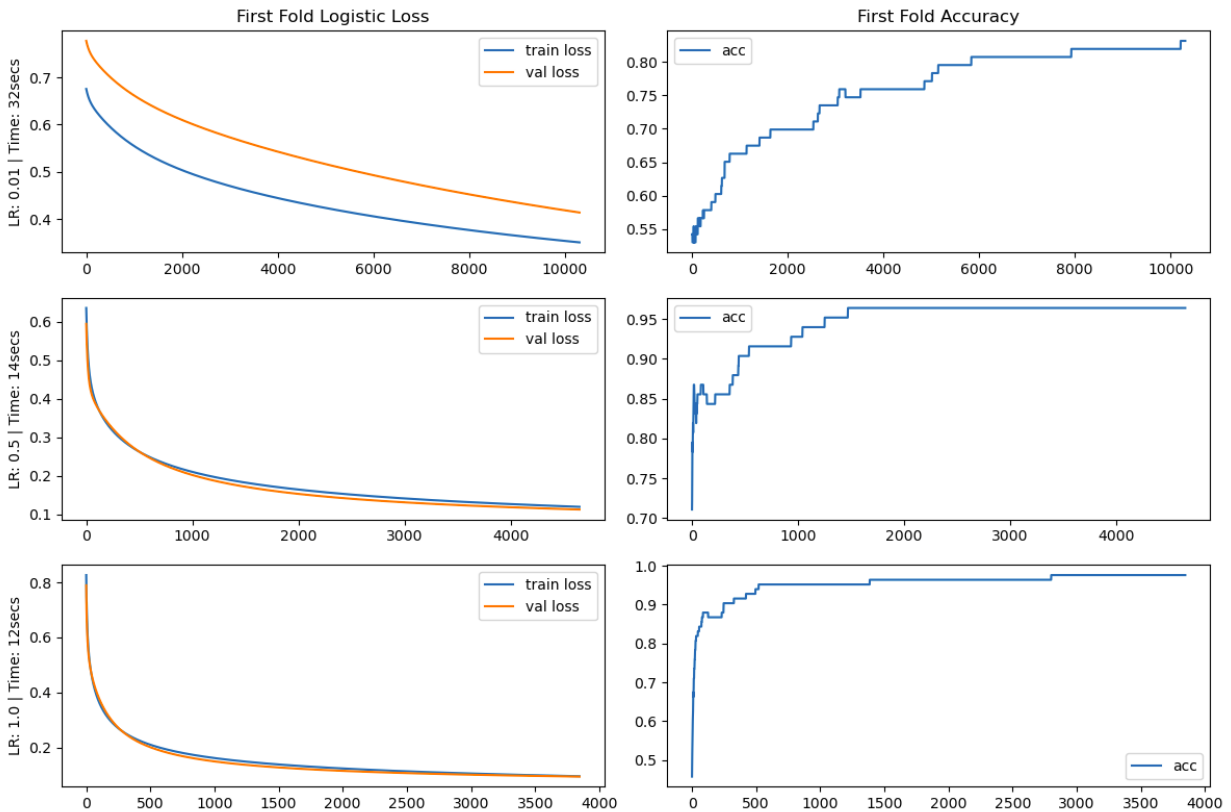


Figure1 Loss and Accuracy values of the first fold when full gradient batch is used.

Increasing the learning rate makes convergence faster in terms of both time and number of iterations. The model did not diverge when LR is 1.0 since we normalize the data.

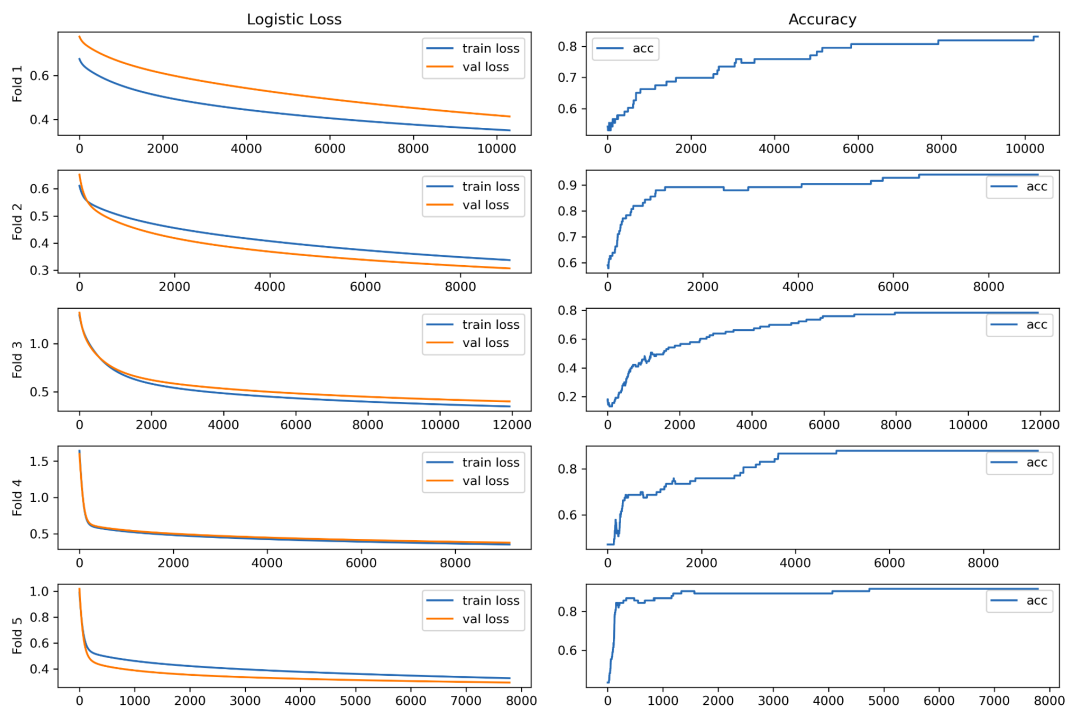


Figure2 Full Gradient Batch & LR:0.01

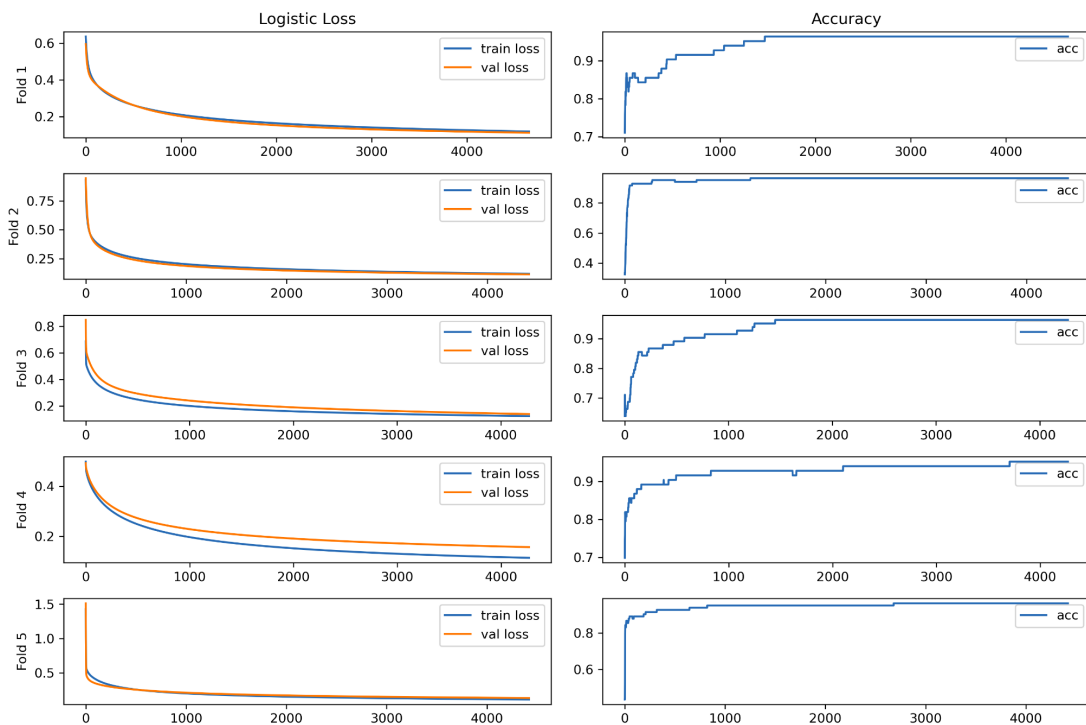


Figure3 Full Gradient Batch & LR:0.5

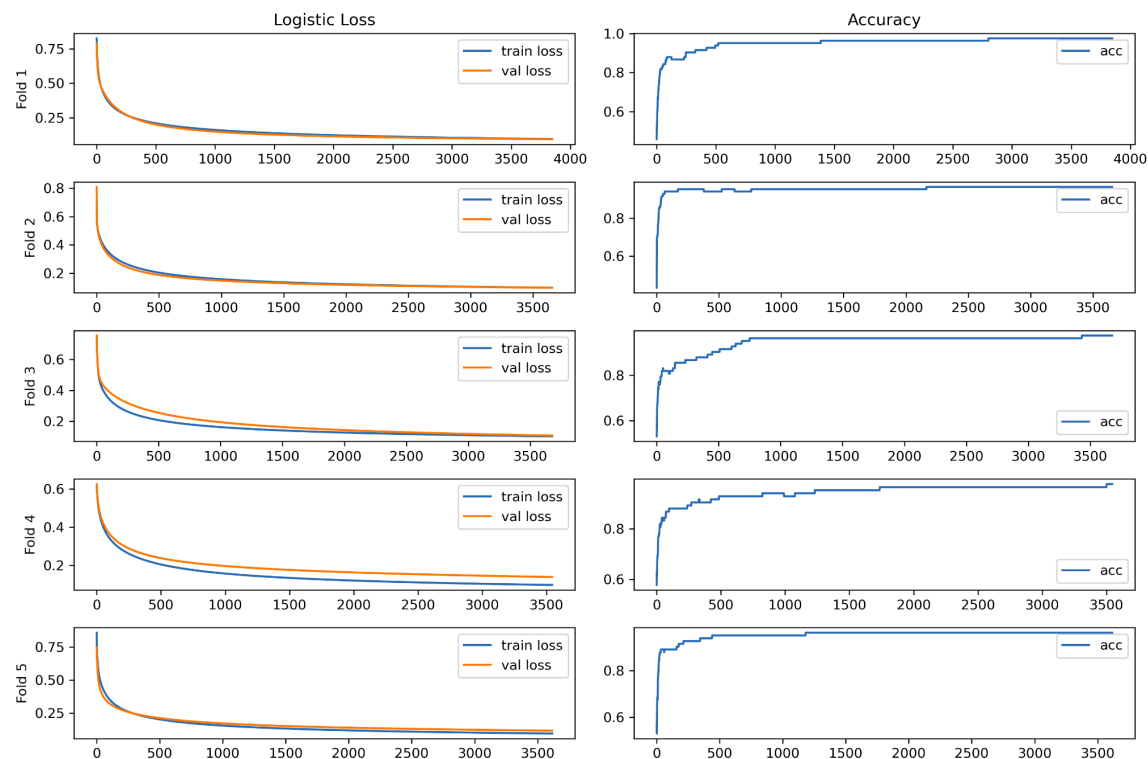


Figure4 Full Gradient Batch & LR:1.0

By default only the **Figure 1** is created after the run. In order to obtain **Figure2&3&4** please make the debug variable **True** from the code.

The acceleration effect of the larger learning rate can also be seen from the **Figure2&3&4**.

Step 2

For the step2 I've used **64 as the mini batch size**.

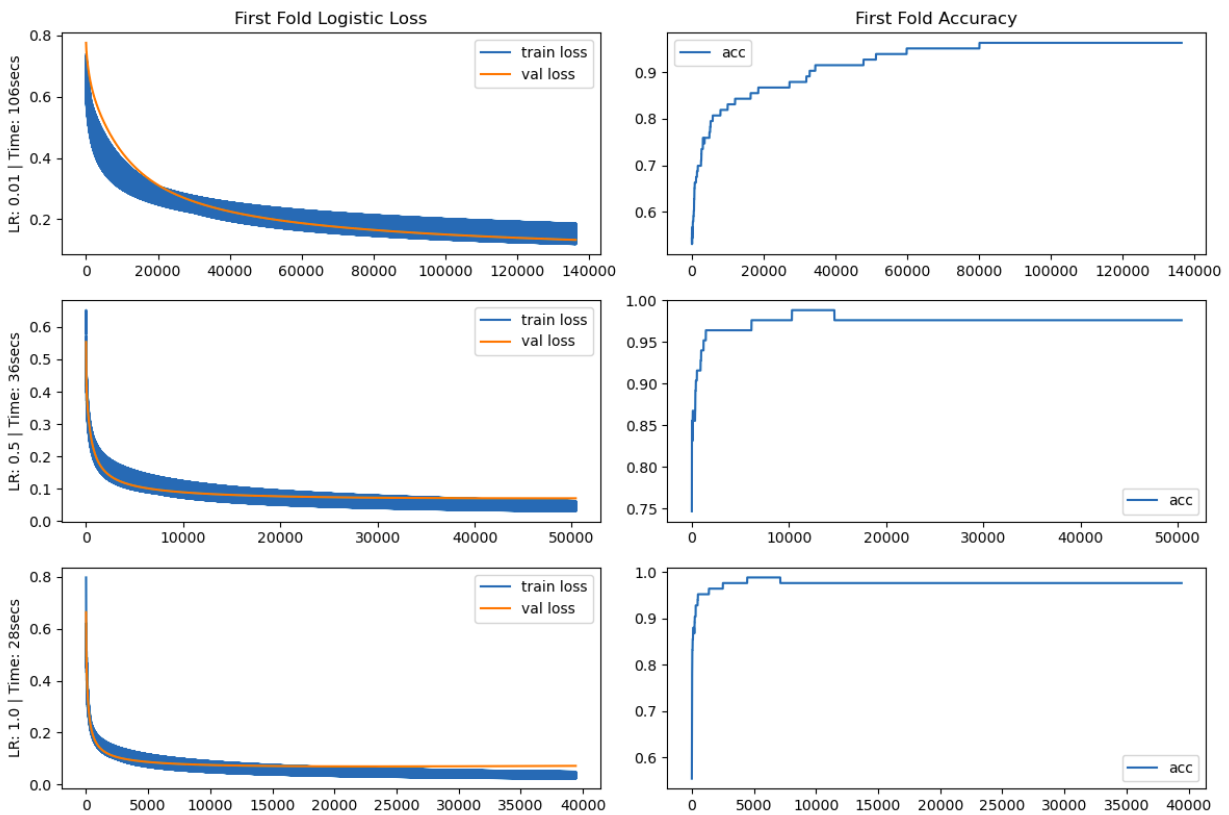


Figure 5 The logistic loss and the accuracy graphs when the Mini Batch Gradient is used.

The training loss value so fluctuates that its graph looks like a region. This fluctuation makes the convergence harder but it reaches a very high accuracy and a very low loss value immediately. However, for a long time, it fluctuates around the optimal weights. Thus, the convergence criteria is met later.

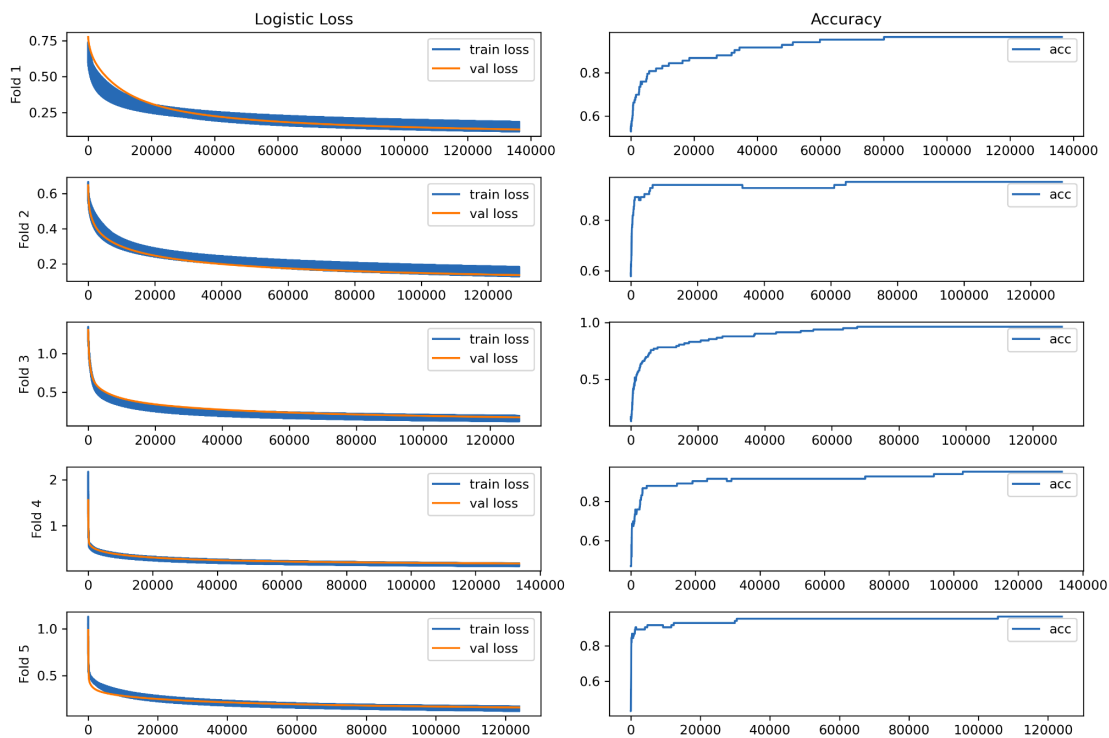


Figure 6 Mini Batch Gradient Descent & LR:0.01

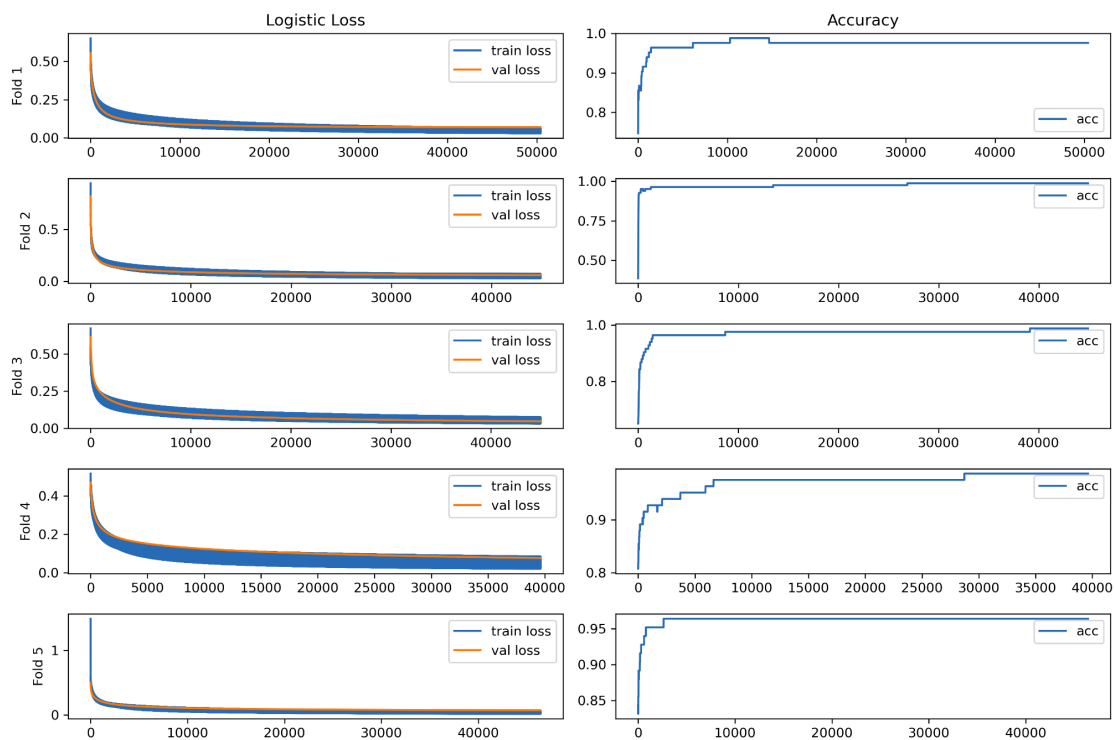


Figure 7 Mini Batch Gradient Descent & LR:0.5

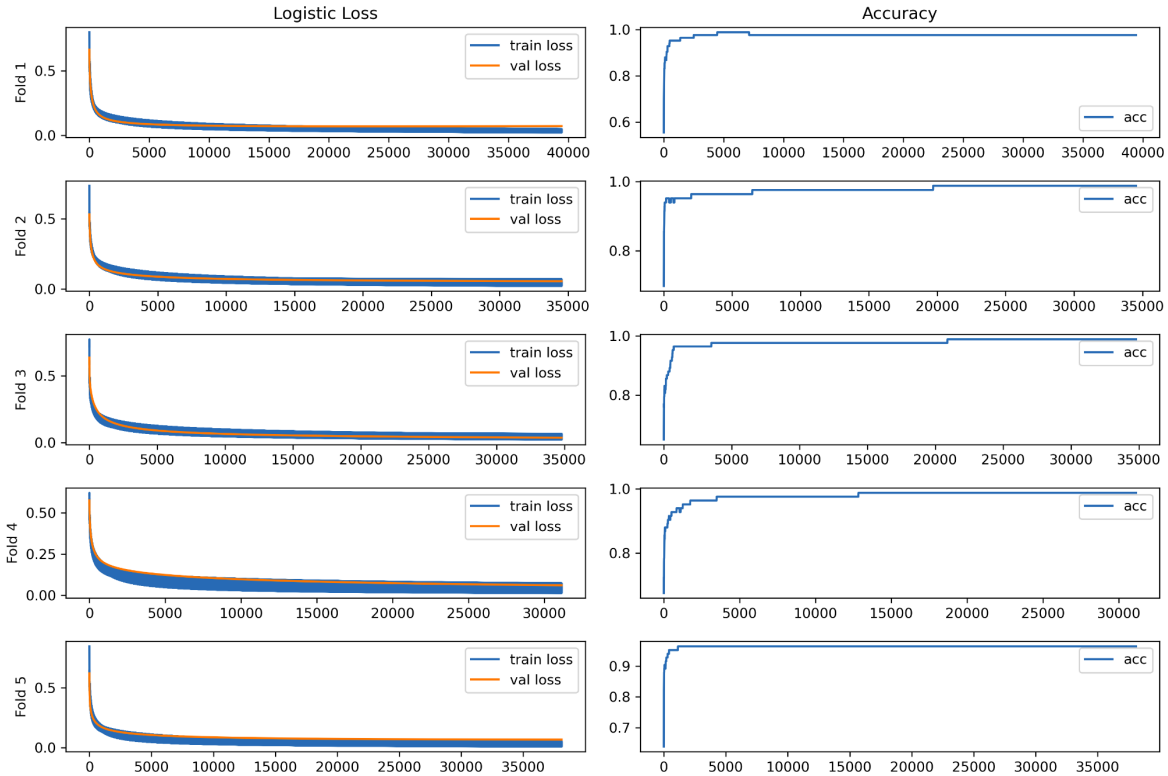


Figure 8 Mini Batch Gradient Descent & LR:1.0

By default only the **Figure 5** is created after the run. In order to obtain **Figure6&7&8** please make the debug variable **True** from the code.

Part 2

In naive bayes, we try to estimate and compare the probability of a sample X belonging to a class y_k . $P(Y = y_k|X)$ In order to obtain the probability of the sample X belongs to each class, we use the statistics we collected from the training set. Our prediction y^* for a sample X can be formulated like this:

$$y^* = \operatorname{argmax}_{y_k} p(Y = y_k|X)$$

If we apply the Bayes Rule we obtain:

$$y^* = \operatorname{argmax}_{y_k} \frac{p(X|Y = y_k)p(Y = y_k)}{p(X)}$$

If we treat every feature as independent and identically distributed, we can rewrite the formula like this:

$$y^* = \operatorname{argmax}_{y_k} \prod_j p_j(x^j|Y = y_k)p(Y = y_k)$$

This assumption makes out Bayes classifier 'Naive'.

For the problem, the probability of being 'mammal' of our test sample is:

$$\begin{aligned}
p(\text{Class} = \text{mammals} | X = \text{test}) = & \\
& p(\text{GiveBirth} = \text{yes} | \text{Class} = \text{mammals}) * \\
& p(\text{CanFly} = \text{no} | \text{Class} = \text{mammals}) * \\
& p(\text{LiveInWater} = \text{yes} | \text{Class} = \text{mammals}) * \\
& p(\text{HaveLegs} = \text{no} | \text{Class} = \text{mammals}) * \\
& p(\text{Class} = \text{mammals})
\end{aligned}$$

$$\begin{aligned}
p(\text{Class} = \text{mammals} | X = \text{test}) &= \frac{6}{7} * \frac{6}{7} * \frac{2}{7} * \frac{2}{7} * \frac{7}{20} \\
p(\text{Class} = \text{mammals} | X = \text{test}) &= \frac{1008}{48020}
\end{aligned}$$

The probability of being 'non-mammal' is this:

$$\begin{aligned}
p(\text{Class} = \text{non} - \text{mammals} | X = \text{test}) = & \\
& p(\text{GiveBirth} = \text{yes} | \text{Class} = \text{non} - \text{mammals}) * \\
& p(\text{CanFly} = \text{no} | \text{Class} = \text{non} - \text{mammals}) * \\
& p(\text{LiveInWater} = \text{yes} | \text{Class} = \text{non} - \text{mammals}) * \\
& p(\text{HaveLegs} = \text{no} | \text{Class} = \text{non} - \text{mammals}) * \\
& p(\text{Class} = \text{non} - \text{mammals})
\end{aligned}$$

$$\begin{aligned}
p(\text{Class} = \text{non} - \text{mammals} | X = \text{test}) &= \frac{1}{13} * \frac{10}{13} * \frac{3}{13} * \frac{4}{13} * \frac{13}{20} \\
p(\text{Class} = \text{non} - \text{mammals} | X = \text{test}) &= \frac{1560}{571220}
\end{aligned}$$

Since $P(\text{Class} = \text{mammals} | X = \text{test})$ is greater than $P(\text{Class} = \text{non} - \text{mammals} | X = \text{test})$, we make our prediction as the test sample X belong to the class 'mammal'.